

Fraction of End-to-End Latency Spent in Isolated Model Execution

